

Historical harvest data - importing it, and what the numbers mean

Both sections below were in Boord's `MANUAL.md` while the historical tables and the Historical Harvest Data report lived there. They describe the import scripts in `owner-app/scripts/` and the one real caveat in the figures, so they belong with whatever rebuilds that report.

Paths are as they were on the farm server (`C:\Boord`), and will need adjusting for wherever this app ends up.

Re-importing historical data on a server

The source workbooks are not shipped with the app. They hold one farm's own harvest records, so they live in `data\imports\` - which is gitignored and per-farm - rather than in the repository. A new install has no workbooks and the two harvest import scripts simply skip, which is correct: a farm should never be importing another farm's history. To load your own, put the workbook in `data\imports\` (or pass its path as the first argument to the script) and re-run.

The historical import is a one-off script, not something an update carries over - the source workbook and the script that reads it are both in the repo, but running the script is a separate manual step against *that server's own database*. This means a fresh install, or a server whose database was reset, needs the import run on it directly - a pull alone leaves the Historical Harvest Data report with nothing before the current season, until this is done.

From Command Prompt in `C:\Boord`, after confirming the update (`update_server.bat` - see [Getting the code onto the server](#)) has already brought in the latest code:

```
backend\.venv\Scripts\python.exe scripts\import_historical_harvest.py
```

Safe to re-run any time (e.g. after regenerating the source workbook) - it replaces the whole historical table each time rather than appending. No server restart needed; the next report picks it up.

There's a second, separate import for the even-older 1987-2019 seasons (annual totals only, no daily breakdown - these feed only the Historical Harvest Data report's Annual Totals sheet). 2012-2019 has a per-block breakdown; 1987-2009 only has a whole-farm total per year, since those records predate today's block register and use an incompatible numbering scheme.

`update_server.bat` runs this automatically on every update (see "Pulling future updates" above), so it's rarely needed by hand - but for a fresh install, or to run it in isolation:

```
backend\.venv\Scripts\python.exe scripts\import_historical_annual_yield.py
```

Safe to re-run any time; replaces its table wholesale.

A note on the historical numbers

A handful of today's blocks (**8a/8b**, **10a/10b**, **17a/17b**, **19a/19b**) didn't exist as separate blocks before the app - the original spreadsheets recorded one combined daily total for each pair. Those combined totals have been split between the two sub-blocks in proportion to their hectares (e.g. a day's picking on the old "block 8" is split roughly 58%/42% between 8a and 8b, matching their relative size). This is a reasonable **estimate**, not what was actually picked from each sub-block on that day - every figure built from an estimated split carries a small info icon next to its block name in the Per-Block Yield table so it's never mistaken for an exact historical record.

When Boord's block ids change

The source workbooks name their columns with the block ids that were current when they were prepared. Boord's register moves — blocks get split, merged, relettered — and a historical row keyed to an id that no longer exists is the **quiet** failure of this whole import. Nothing errors. Season totals still add up, because the Analysis tab sums by season as well as by block. The block simply shows no history, and a block with no history looks exactly like a block that was not bearing yet.

It has happened once already. Three pairs lost their **a / b** lettering, with the **second** of each pair taking the bare number:

In the workbooks	In Boord today	
10a Mauritius, 512 trees, 2.5 ha	10b Block 10b-Mau	
10b Early Delight, 512 trees, 2.4 ha	10 Block 10-ED	
17a Mauritius, 584 trees, 2.8 ha	17b Block 17b-Mau	
17b Early Delight, 584 trees, 2.8 ha	17 Block 17-ED	
19a Mauritius, 358 trees, 1.7 ha	19b Block 19b-Mau	
19b Early Delight, 358 trees, 1.7 ha	19 Block 19-ED	

Note this is the **opposite of what the ids suggest**: the bare **10** is the old **10b**, not the old **10a**. The mapping was established from variety, tree count and hectares, which agree on all six — hectares settle it, since new **10** is 2.4 ha and old **10a** was 2.5. Mapping by id instead would put each Mauritius block's six seasons onto an Early Delight block of a different size; season totals would still be right, and only kg/ha, kg/tree and the variety charts would be wrong, against a past nobody can check by eye.

8a / 8b were not affected.

The map lives in `scripts/block_renames.py` and both import scripts apply it. **Add to it rather than editing the workbooks** — those are the farm's own records and should keep saying what they said when they were written.

After any import, and after any Boord release that touches Master Data:

```
backend\.venv\Scripts\python.exe scripts\check_block_ids.py
```

It lists any imported id Boord no longer has, and exits non-zero — so a future rename is reported out loud instead of silently emptying a block.